

WHITE PAPER COMPANION · I

SILICONA

Competitor Intelligence

Evidence-Scored Landscape for the Engineering Control Plane — v1.0. Fourteen primary sources fetched live, every competitor claim tiered and quoted verbatim, every number attributed — an honest map of who has evidence and who has adjectives.

SILICONA Engineering

Companion to the Technical White Paper — 12 September 2026

Evidence cutoff 12 September 2026 · 14 primary sources

EVIDENCE · MEASUREMENT · VERIFICATION

Contents

1. Method: The Evidence-Tier System	1
Claim-discipline rules	1
What this chapter does not do	2
2. The Landscape: Four Categories	2
3. The Autonomy-Numbers Problem	3
The multiplier wall, quoted verbatim	3
The SILICONA counter-standard, stated once	5
4. Profiles	6
Synopsys	6
Cadence	6
Siemens EDA	7
Ricursive Intelligence	7
ChipAgents	8
Agentrys	9
Cognichip	10
Silimate	10
Rise Design Automation	11
MooresLabAI	11
Agnisys	12
CaretEDA	13
OpenROAD	13
LibreLane	14
5. Funding versus Substance	15
The contrast SILICONA can honestly make	16
6. Positioning: Where the Control Plane Sits	16
Where SILICONA does not compete	16
7. Implications	17
8. Source Register	18

1. Method: The Evidence-Tier System

This document is the competitive-landscape companion to the SILICONA technical white paper. Its subject is the agentic-EDA market as it stood on 12 September 2026: who is funded, who is shipping, what is claimed, and — the part almost nobody publishes — how much of each claim is checkable. SILICONA is an evidence-driven engineering control plane, not “another AI for EDA”, and this chapter is written to that standard: competitors in this market have serious autonomous workflows and well-funded agentic stacks, and nothing in this document argues that SILICONA out-agents them. The argument is narrower and, we believe, more durable: claims compound in value only when they are traceable, and almost none of the claims in this landscape currently are.

Every external fact below carries an evidence tier assigned at fetch time. Four tiers cover the fresh primary sources; a fifth marker covers the small residue of facts that come only from SILICONA’s prior research foundation and were not freshly fetched for this edition.

Tier	Definition	What it can support	Sources in this chapter
[E1]	Independently reported journalism	Funding, investors, headcount, deployment facts — citable with publisher attribution; quoted claims remain vendor statements to press	Reuters (ChipAgents), TechCrunch (Rrecursive)
[E2]	Vendor press release	Roadmap, partnership, product-naming and availability facts; every performance multiplier stays company-reported	Synopsys ×2, Cadence, Agentrys, CaretEDA, Agnisys
[E3]	Vendor self-published product page	Positioning language only; numeric claims are marketing until independently corroborated	silimate.com, mooreslab.ai, rise-da.com
[E4]	Open-source verifiable artifact	The strongest transparency tier: claims inspectable in repos, docs and shuttles	openroad.org, GitHub/OpenROAD, librelane docs
[S]	Secondary foundation research (not freshly fetched for this edition)	Context only; flagged wherever used, never load-bearing for a SILICONA contrast	Siemens EDA, Cognichip, market-economics context

Table 1: The evidence tiers. A fact’s tier limits what it may be used for.

Claim-discipline rules

- **Quotes are verbatim.** Competitor claims appear in quotation marks exactly as published; tier tags attach to every fact.
- **Company-reported numbers stay company-reported.** No multiplier, percentage or speedup in this chapter is endorsed, repeated as fact, or averaged with measured data.
- **Pre-GA availability is flagged.** Capabilities announced for “second half of 2026” availability are labelled pre-GA throughout.
- **Distributions beat headlines.** Where a competitor publishes a single headline multiplier with no protocol, sample size, seeds or variance, this chapter says so — and points at the benchmark-protocol companion that defines the counter-standard.

- **SILICONA’s own numbers are DB-traceable.** Measured figures come from the engineering-memory database; the claimed capability ceiling is “demonstrated open PDK closure” (sky130A: DRC 0, LVS MATCH dual verdict, GDS-II streamout). Post-route extraction and signoff STA are explicitly unclaimed. When a competitor’s “sign-off” language is quoted, it is their word, in quotation marks, not ours.

What this chapter does not do

- It makes **no superiority claims** — not one sentence argues SILICONA beats any named competitor on any axis.
- It does **not compare iteration speed against human engineers** (a standing do-not-publish rule from the white paper claim ledger).
- It does **not claim production readiness** for SILICONA — the platform’s own published gaps stay published.
- It does not present competitor marketing numbers as measurements; the tier tag travels with every figure, including in tables.

14	4	14	0
primary sources fetched live	evidence tiers applied	entities profiled	superiority claims made

2. The Landscape: Four Categories

The agentic-EDA market of September 2026 splits into four structural categories. Big-EDA platforms (Synopsys, Cadence, Siemens EDA) bolt agentic layers onto trusted commercial engines. Funded agentic startups (Ricursive, ChipAgents, Agentrys, Cognichip, Silimate, Rise DA, MooresLabAI, Agnisys) sell autonomous workflow products at venture speed. A commercial open-stack player (CaretEDA) productizes the same open toolbase SILICONA builds on. And the open-source flows (OpenROAD, LibreLane) are the verifiable substrate under several of the commercial stories. Table 2 anchors each entity with one tiered fact; Chapter 4 profiles them individually.

Entity	Category	Anchor fact (tier-tagged)	Tier
Synopsys	Big-EDA platform	AgentEngineer™ autonomous workflows; “up to 50X faster time-to-validated RTL”; AMD “actively evaluating”	[E2]
Cadence	Big-EDA platform	ChipStack™ at self-defined “Level-5” autonomy; “over 40X faster RTL validation”; early access 2H-2026	[E2]
Siemens EDA	Big-EDA platform	EDA AI System positioned as generative/agentic layer; Solido adding agentic AI for custom IC	[S]
Ricursive Intelligence	Funded agentic startup	\$335M raised in ~4 months at reported \$4B valuation; Nvidia an investor; customers unnamed	[E1]
ChipAgents	Funded agentic startup	\$131M Series A total; B Capital lead; Micron, MediaTek, Ericsson prior investors; ~64 employees	[E1]

Entity	Category	Anchor fact (tier-tagged)	Tier
Agentrys	Funded agentic startup	\$24.5M total; claims 32-bit CPU “specification to sign-off-clean GDS” and NVIDIA CVDP “exceed 90% accuracy”	[E2]
Cognichip	Funded agentic startup	“Artificial Chip Intelligence” on a physics-informed foundation model; \$60M Series A, \$93M+ total (April 2026)	[S]
Silimate	Funded agentic startup	“11x faster bug resolution” / “70x faster PPA optimization” from anonymous testimonials	[E3]
Rise Design Automation	Funded agentic startup	“Propose, Evaluate, Decide flow grounded in real tool metrics”; “verification at 30-100x RTL speeds”	[E3]
MooreLabAI	Funded agentic startup	“Deployed at 15+ chip companies”; FAQ ROI “92–97% productivity gains”	[E3]
Agnisys	Spec-automation vendor	IDS-AI™ agentic spec framework; “4× improvement in TCPF and 50% fewer top-level ATPG patterns” (Google-co-authored talk)	[E2]
CaretEDA	Commercial open-stack	DAC 2026: commercially available end-to-end open EDA stack (Verilator/Yosys/SymbiYosys/OpenROAD) + 90-day startup program	[E2]
OpenROAD	Open-source flow	“over 600 silicon-ready tapeouts” up to 12nm; 24-hour NHIL flow; 3.1k stars, BSD-3	[E4]
LibreLane	Open-source flow	OpenLane successor under FOSSi Foundation stewardship after the Efabless liquidation	[E4]

Table 2: The landscape at a glance — category, anchor fact, tier.

Two structural observations frame everything that follows. First, incumbents and startups now speak the same language — “fully autonomous”, “no human in the loop”, “Level-5” — which makes autonomy claims undifferentiated as marketing and worthless as evidence. Second, the only checkable numbers in the landscape come from outside its commercial layer: OpenROAD’s inspectable tapeout record [E4], Agentrys’ benchmark-attached CVDP number [E2]. The scarce asset in this market is not autonomy; it is provability.

3. The Autonomy-Numbers Problem

Every serious player in this market publishes dramatic performance multipliers. None of them — not one of the fourteen sources fetched for this chapter — publishes the protocol behind the number: no sample size, no seeds, no variance, no environment pinning, no per-run provenance. The contention here is not that the numbers are wrong; it is that as published they are unfalsifiable, and unfalsifiable numbers cannot compound into engineering knowledge.

The multiplier wall, quoted verbatim

“Each engineer will use ChipStack agents to run hundreds of dynamic simulations with Cadence® Xcelium™ Logic Simulation and Jasper® Formal Verification, delivering over 40X faster RTL validation cycles and reducing a typical five-week verification loop to less than a day.” — Cadence press release, 1 June 2026 [E2; early access planned for 2H-2026]

Cadence’s “Level-5” is the company’s own autonomy taxonomy, not an external standard, and the release itself scopes availability: capabilities are “expected to be available to early-access customers in the second half of 2026” [E2] — pre-GA at fetch time. The release names the engines (Xcelium, Jasper) and the partner stack (NVIDIA Nemotron models, OpenShell runtime), which is genuinely more specificity than most competitors offer; what it does not name is any measurement basis for the 40X or the five-weeks-to-a-day claim.

“The end-to-end fully autonomous verification closure agentic flow compresses weeks of manual labor into hours of agentic execution that achieves up to 50X faster time-to-validated RTL with an additional 20% improvement in coverage.” — Synopsys press release, 26 July 2026 [E2; availability planned for 2H-2026]

Synopsys’s companion release (27 July 2026, the AMD/Microsoft Discovery announcement) carries a detail worth quoting precisely because it is the norm rather than the exception: the headline states “initial results demonstrate up to 40% reduction in cycle time” while the body softens to “Early evaluations show reductions of 25–40% in debug cycle time, saving many weeks of engineering efforts” [E2]. Both figures are Synopsys’s; we quote both and note the discrepancy honestly rather than pick one. The same release reports customer status in two honest words: AMD is “actively evaluating” [E2] — evaluating, not deployed.

The startups mirror the pattern with even less scaffolding. Silimate repeats “11x faster bug resolution” and “70x faster PPA optimization” four times across its homepage, attributed to anonymous testimonials (“VP of Engineering”, “Senior Principal Engineer”) [E3]. MooresLabAI’s FAQ publishes “92–97% productivity gains” and “200+ engineering hours saved per project” as self-reported ROI, alongside “VerifAgent™ reduces verification costs by up to 86%. Manual UVM verification that typically takes 6–12 months can now be completed in under a month.” [E3]. Rise DA claims “verification at 30–100x RTL speeds” [E3]. Agnisys is the most specific — “Early adopters have seen up to 4× improvement in Test Coverage Per Flip-Flop (TCPF) and 50% fewer top-level ATPG patterns” — presented in a DAC 2026 talk co-authored with Google engineers, though the early adopters remain unnamed [E2].

Vendor	Headline claim (verbatim, tier)	Published protocol?	What is missing
Cadence	“over 40X faster RTL validation”; “five-week verification loop to less than a day” [E2]	None published	n, methodology, environment; pre-GA (2H-2026 early access); “Level-5” is a self-defined taxonomy
Synopsys	“up to 50X faster time-to-validated RTL” + “20% improvement in coverage”; debug “25–40%” (headline “up to 40%”) [E2]	None published (footnote markers unresolved)	n, seeds, variance; 2H-2026 availability; headline/body inconsistency unexplained
Silimate	“11x faster bug resolution” / “70x faster PPA optimization” [E3]	None	Attribution (anonymous testimonials), methodology, designs, environment
Moore'sLabAI	“92–97% productivity gains”; “reduces verification costs by up to 86%” [E3]	None	Sample, definition of “productivity”, customers, baseline
Rise DA	“verification at 30-100x RTL speeds” [E3]	None	Meaning of “x RTL speeds” (virtual-prototyping framing), methodology
Agnisys	“up to 4x improvement in TCPF and 50% fewer top-level ATPG patterns” [E2]	Conference talk, no protocol page	Early-adopter identities, n, variance; numbers vendor-reported until paper

Table 3: The autonomy numbers, as published — and what is missing behind each.

The SILICONA counter-standard, stated once

SILICONA’s answer to the multiplier wall is not a bigger multiplier; it is a published protocol under which any SILICONA number can be re-derived. Its companion document — the SILICONA Benchmark Protocol (v1.0, this edition) — requires $n \geq 5$ repetitions per family per mode, recorded seeds, fresh-context AI episodes, fresh re-measurement of winners (citation never substitutes for measurement), pinned tool versions, failure rows published as data (including the 51 honest HTTP-429 provider failures already in the ledger), mandatory negative results (a failed search family is published as failed), and distributions (min/median/max/IQR plus the full per-run table) instead of single headline percentages. The house rule reads: AI proposes; deterministic EDA tools measure; verification decides; every result recorded as traceable engineering evidence.

4. Profiles

Fourteen profiles follow, each built the same way: a facts table, a split between what is demonstrated (as evidenced at the tier cited) and what is asserted (company claim, unverified), and a short analysis ending in what the profile means for an evidence-control-plane thesis. Quotation marks are reserved for the competitors’ own words.

Synopsys

Category	Evidence tier	Funding anchor	Availability
Big-EDA platform	Press release ×2 [E2]	Public company (none)	Evaluation now; GA “planned for the second half of 2026” [E2]

Profile: Synopsys, Inc. (NASDAQ: SNPS) — two merged releases, 26–27 July 2026

Demonstrated (as evidenced)	Asserted (company claim, unverified)
Two named workflow families on named infrastructure: AgentEngine™ on Microsoft Discovery, NVIDIA Nemotron 3 Ultra + OpenShell, Fusion Compiler on Azure [E2]	“fully autonomous”, “long-running” agent behavior; “up to 50X faster time-to-validated RTL”; “20% additional coverage improvement” [E2]
Named evaluators: AMD (“actively evaluating”), Microsoft as platform partner; analog (up to 3X) and CAE workflows demoed at DAC 2026 [E2]	Debug-closure gains: headline “up to 40% reduction in cycle time” vs body “25–40%” — both quoted; inconsistency unexplained [E2]
Breadth: 20+ GPU-enabled EDA/multiphysics products; PrimeSim “up to 18X”, QuantumATK “up to 50X / 200x”, Lumerical “10X” (all vendor-reported) [E2]	“Industry’s First Autonomous EDA Workflows on Microsoft Discovery” — superlative scoped to the platform [E2]

Synopsys — demonstrated vs asserted

Synopsys is the landscape’s most credible agentic incumbent on partnerships and the least careful on number hygiene: the 50X/20% claims ride unresolved footnote markers, and the same week’s two releases disagree on whether debug closure is “up to 40%” or “25–40%”. None of this makes the claims false; it makes them unverifiable, which for an engineering buyer is the same problem. For an evidence-control-plane thesis, Synopsys defines the incumbent pattern to complement, not beat.

Cadence

Category	Evidence tier	Funding anchor	Availability
Big-EDA platform	Press release [E2]	Public company; ChipStack acquired Nov 2025 [E2]	Pre-GA: “early-access customers in the second half of 2026” [E2]

Profile: Cadence Design Systems (NASDAQ: CDNS) — Computex 2026 release, 1 June 2026

Demonstrated (as evidenced)	Asserted (company claim, unverified)
Portfolio timeline on record: ChipStack acquired Nov 2025, first product Feb 2026, agent portfolio at CadenceLIVE April 2026 [E2]	“industry’s first fully autonomous virtual agentic AI design engineer” — superlative claim [E2]
Named engines and runtime: Xcelium, Jasper, NVIDIA Nemotron models, OpenShell sandbox; ViraStack/InnoStack/AgentStack portfolio named [E2]	“Level-5 autonomy” — Cadence’s own taxonomy, not an external standard [E2]
Named differentiator claim: autonomous behavior “tightly coupled” with “signoff-accurate” core engines [E2]	“over 40X faster RTL validation cycles”; “reducing a typical five-week verification loop to less than a day” — no methodology [E2]

Cadence — demonstrated vs asserted

Cadence’s release is the most quotable in the landscape precisely because it commits to specific imagery — Level-5 autonomy, hundreds of dynamic simulations, five weeks collapsing to a day — while shipping none of the measurement context that would let an engineer check any of it. The genuinely strong part is structural: agents coupled to Xcelium and Jasper mean the deterministic-measurement half of an evidence loop already exists inside their stack; what is missing is the publication layer. An evidence control plane is that missing layer, not a missing engine.

Siemens EDA

Category	Evidence tier	Funding anchor	Availability
Big-EDA platform	Secondary foundation research [S] — flagged	Public company (Siemens AG)	Not freshly verified; no primary fetch in this pass

Profile: Siemens EDA — foundation research only; NOT freshly fetched for this edition (secondary tier, flagged)

Demonstrated (as evidenced)	Asserted (company claim, unverified)
Portfolio positioning on record via prior research: “EDA AI System as a generative and agentic layer across its EDA portfolio, with centralized multimodal design data and RAG” [S]	“Solido is adding agentic AI capabilities for custom IC design, simulation, verification, and optimization.” [S]
None at primary tier — no fresh fetch; treat all Siemens facts as context until a primary source is on file	Strategic framing from the foundation: analog/custom design as a major future battleground [S]

Siemens EDA — demonstrated vs asserted (all secondary-tier)

Siemens EDA appears in this chapter under an explicit handicap: no primary source was fetched for this edition, so every Siemens fact carries the [S] secondary tag and nothing more. The foundation’s reading — an agentic layer across the portfolio plus Solido agentic AI for custom IC — is plausible, but plausibility is not evidence, and this profile is written to model that distinction. A follow-up primary fetch is on file as a to-do before any stronger statement is made.

Ricursive Intelligence

Category	Evidence tier	Funding anchor	Availability
Funded agentic startup (AI-first chip designer)	Journalism [E1] — TechCrunch, Julie Bort	“\$300 million Series A round at a \$4 billion valuation led by Lightspeed” months after “a \$35 million seed round led by Sequoia” [E1]	No shipped product detailed in coverage

Profile: Ricsursive Intelligence — TechCrunch, 16 February 2026

Demonstrated (as evidenced)	Asserted (company claim, unverified)
Funding facts independently reported: \$335M in ~4 months, \$4B valuation, Lightspeed and Sequoia leads, Nvidia an investor (“Nvidia is an investor”) [E1]	Platform will “learn across different chips” — each design improving the next [E1, founder quote]
Founder pedigree on record: Anna Goldie and Azalia Mirhoseini, AlphaChip lineage at Google, three TPU generations [E1]	“We could design a computer architecture that’s uniquely suited to that model, and we could achieve almost a 10x improvement in performance per total cost of ownership” — explicitly hypothetical [E1]
Zero shipped-product evidence in the journalism: no named customers, no product details, no benchmarks, no tapeouts [E1]	Customer traction narrative: founders say they have “their pick of their first development partners” — customers unnamed [E1]

Ricsursive — demonstrated vs asserted

Ricsursive is the clearest case in the landscape of money preceding evidence: \$335M raised in four months on the strength of a research pedigree, with the journalism itself recording zero shipped-product detail and unnamed early customers. Nothing here diminishes the pedigree — AlphaChip is real published research, and the founder lineage is journalism-grade [E1] evidence — but the product story remains founder statements to press, which is a different evidence class entirely. For an evidence-control-plane thesis, Ricsursive marks the market’s appetite: investors are paying now for learning-across-designs stories whose verification machinery does not yet exist in public — which makes the verification machinery itself the complementary build.

ChipAgents

Category	Evidence tier	Funding anchor	Availability
Funded agentic startup (verification-first)	Journalism [E1] — Reuters, Max A. Cherney	“B Capital led the additional funding round, which brings its Series A fundraising total to \$131 million.” Prior investors: Micron, MediaTek, Ericsson [E1]	Product sold; deployments not detailed in coverage

Profile: ChipAgents — Reuters exclusive, 29 July 2026

Demonstrated (as evidenced)	Asserted (company claim, unverified)
Funding facts independently reported: \$60M expansion, \$131M Series A total, B Capital lead [E1]	“ChipAgents sees the largest speedup in design with the company’s software that is designed to help verify a chip will function as intended” — CEO to Reuters; unquantified [E1]
Company facts: ~64 employees, Santa Clara HQ; expanded Nvidia collaboration on a specialized chip-design AI model; investor status of Nvidia declined to comment [E1]	Agent capability framing: programs “capable of making decisions and executing complex tasks with little or no human oversight” [E1]
Named strategic investors from prior rounds: Micron, MediaTek, Ericsson [E1]	No customer names, benchmarks or tapeouts in the coverage [E1]

ChipAgents — demonstrated vs asserted

ChipAgents is the best-funded company in the landscape whose headline speedup claim is, honestly, no number at all: the CEO’s on-record position is that verification is where the software helps most — an assertion Reuters reports without quantification. The verifiable facts are the capital and the investor list (strategic money usually means pilots). For the evidence-control-plane thesis, ChipAgents validates that verification is the wedge — and leaves the measured-verification-evidence position unoccupied.

Agentrys

Category	Evidence tier	Funding anchor	Availability
Funded agentic startup (open-platform positioning)	Press release [E2]	“\$24.5 million... an oversubscribed \$19.1 million seed round led by Etna Labs, following a \$5.4 million pre-seed led by MediaTek” [E2]	Platform claimed demonstrated end to end (vendor-reported)

Profile: Agentrys — vendor announcement, 26 August 2026

Demonstrated (as evidenced)	Asserted (company claim, unverified)
Concrete funding structure with named leads (Etna Labs; MediaTek pre-seed) [E2]	Full-flow demo: an autonomous multi-agent workflow “took a 32-bit CPU from specification to sign-off-clean GDS layout with no human in the loop” — vendor-reported; “sign-off-clean” is Agentrys’s word [E2]
The only named external AI benchmark in the landscape with a number: “On NVIDIA’s public CVDP verification benchmark, Agentrys is among the first to exceed 90% accuracy” [E2] — checkable in principle against the public benchmark	Customers: “several of the world’s top fabless semiconductor companies, a leading global foundry, and several leading chip startups” — all unnamed [E2]
Team lineage on record: NVIDIA/Meta/AMD/Samsung/Google/Siemens researchers; ChipNeMo, NVCell, DREAMPlace, VerilogEval, CVDP credits [E2]	Category-defining language: “Agentic Design Automation (ADA)” evolving EDA “from traditional EDA tools focused on individual tasks to intelligent systems” [E2]

Agentrys — demonstrated vs asserted

Agentrys matters to this chapter for one structural reason: it is the only vendor in the fourteen sources that attaches a number to a named public benchmark — NVIDIA CVDP, “exceed 90% accuracy” — which makes the claim falsifiable even though the release is vendor-tier [E2]. Its autonomy claims (“no human in the loop”, “sign-off-clean” — their words, quoted) are exactly the kind SILICONA does not make; SILICONA’s equivalent is the narrower, tool-verifiable “demonstrated open PDK closure”. That the industry reads a benchmark-attached number as extraordinary is the strongest available evidence that the counter-standard is needed.

Cognichip

Category	Evidence tier	Funding anchor	Availability
Funded agentic startup (AI-first chip designer)	Secondary foundation research [S] — flagged	\$60M Series A, “more than \$93 million total” (April 2026) [S]	Not freshly verified; no primary fetch in this pass

Profile: Cognichip — foundation research only; NOT freshly fetched (secondary tier, flagged)

Demonstrated (as evidenced)	Asserted (company claim, unverified)
Funding facts as recorded by prior research: \$60M Series A announced April 2026; \$93M+ total [S]	“Artificial Chip Intelligence”, based on “a physics-informed foundation model for chip design” [S]
None at primary tier — no fresh fetch; all facts carry the secondary flag	No measured claims available at any tier in this chapter’s evidence base

Cognichip — demonstrated vs asserted (all secondary-tier)

Cognichip shares Rrecursive’s archetype — an AI-first chip designer with a foundation-model story and nine-figure backing — and shares this chapter’s evidence handicap: no fresh primary source, so the [S] flag rides every fact, including the \$93M+ total. The profile exists for completeness of the funded cohort, not for any claim it can support: the second AI-first designer funded at scale on model-prowess narrative, with nothing in the public record a measurement standard could grip. A primary fetch is on the to-do list before the next edition.

Silimate

Category	Evidence tier	Funding anchor	Availability
Funded agentic startup (frontend/PPA + debug)	Vendor marketing [E3]	No primary funding figures on page; logo strip only	Commercial product (licensed); no named customers

Profile: Silimate — company homepage, live © 2026

Demonstrated (as evidenced)	Asserted (company claim, unverified)
Named public posture: on-prem/self-hosted, “No egress of your data at any point” [E3]	“11x faster bug resolution” / “70x faster PPA optimization” — repeated 4x on the page, attributed to anonymous testimonials [E3]
Marketing design attribution: “Design shown is from the OpenTitan project” — the one named design is an open project, not a customer [E3]	“Converge to your PPA targets and fix functional bugs in days instead of months”; “runs on highly-complex designs with millions of gates” [E3]
Their CTO’s OpenROAD testimonial (see that profile) confirms open-tool grounding in their copilots [E4, cross-source]	“Licensed by chip teams building industry-leading SoCs and IPs” — no names [E3]

Silimate — demonstrated vs asserted

Silimate is the landscape’s purest example of the anonymous-testimonial multiplier: two headline numbers, repeated four times, zero named humans behind them, and no funding figures on the page [E3]. The most interesting evidence about Silimate is second-hand: its CTO’s named testimonial on openroad.org praising

OpenSTA and open-source tooling as “physics-accurate grounding” [E4] — a public admission, from a competitor, that open tools are where the grounding lives. Silimate is the marketing end of the spectrum: positioning language for the same loop SILICONA actually instruments.

Rise Design Automation

Category	Evidence tier	Funding anchor	Availability
Funded agentic startup (spec-centric, HLS lineage)	Vendor marketing [E3]	No funding figures on page; Microsoft for Startup Founders Hub affiliation [E3]	Product pages live (IPCreate / IP Transform / High-Level Platform)

Profile: Rise Design Automation — company page, live

Demonstrated (as evidenced)	Asserted (company claim, unverified)
Founder pedigree, verifiable career facts: Badru Agarwala, 40 years in EDA; ex-GM Calypto Systems Division at Mentor (now Siemens EDA); Catapult HLS; two prior companies acquired [E3]	“The Rise AI-Native Platform, driven by a Propose, Evaluate, Decide flow grounded in real tool metrics” [E3]
Product naming on record: Rise AI IPCreate, IP Transform, High-Level platforms [E3]	“verification at 30-100x RTL speeds” [E3]
Zero published evidence of the loop (no runs, no metrics, no protocol)	“the fastest path from Spec to production-quality RTL” — superlative, unfalsifiable as stated [E3]

Rise DA — demonstrated vs asserted

“The Rise AI-Native Platform, driven by a Propose, Evaluate, Decide flow grounded in real tool metrics, delivers what no other approach can: the fastest path from Spec to production-quality RTL, early architectural exploration, early HW/SW validation with RTL accuracy, and verification at 30-100x RTL speeds.” — rise-da.com company page [E3]

Rise DA deserves a longer look than its evidence weight because it is the closest conceptual overlap with SILICONA found across all fourteen sources: a Propose–Evaluate–Decide loop, explicitly “grounded in real tool metrics”, is the same loop an evidence control plane enforces. The difference is the receipts: SILICONA publishes the loop’s outputs (per-run provenance, honest negatives, distributions); Rise DA’s public record contains the phrase and none of them. Convergence from a 40-year EDA veteran is validation; the absent evidence layer is the whole differentiation.

Moore’s Lab AI

Category	Evidence tier	Funding anchor	Availability
Funded agentic startup (verification-first suite)	Vendor marketing [E3]	No investor names or amounts on page; “backed by leading industry partners” [E3]	Products sold (SpecAgent, VerifAgent, DebugAgent, CoverageAgent, SignoffAgent; DesignAgent “coming soon”)

Profile: Moore’s Lab AI — company homepage + FAQ, live © 2026

Demonstrated (as evidenced)	Asserted (company claim, unverified)
Founders named in FAQ: Shelly Henry, Shashank Chaurasia, Sirish Munipalli; founded 2025; “45 years of combined semiconductor experience” [E3]	“Deployed at 15+ chip companies worldwide” [E3]
Named partnership: Breker Verification Systems (“First AI-Driven SoC Verification Solution” co-marketing); Triple Crown partnership [E3]	FAQ ROI: “92–97% productivity gains”, “50% reduction in time to tapeout”, “200+ engineering hours saved per project”, “80% fewer bugs caught in verification” [E3]
Engagement models published (licenses / +FDE / +DV engineers; Verification-as-a-Service “<45 days”) [E3]	“VerifAgent™ reduces verification costs by up to 86%. Manual UVM verification that typically takes 6–12 months can now be completed in under a month.” [E3]

MooresLabAI — demonstrated vs asserted

MooresLabAI is the landscape’s densest self-reported-ROI page: five percentages in one FAQ, a deployment count (“15+ chip companies”) with no names, anonymous testimonials, and a product line where one agent is still “coming soon” [E3]. The Breker partnership is the one externally checkable fact, and it is co-marketing, not measurement. What the profile adds is the shape of the verification-first segment — the same wedge as ChipAgents and Agentrys — occupied at three price points and zero published protocols. For the evidence-control-plane thesis: this is where a measured, per-run-audited alternative would sit, if one existed.

Agnisys

Category	Evidence tier	Funding anchor	Availability
Spec-automation vendor (~2 decades, niche)	Press release + conference talk [E2]	None stated	Products shipping (IDesignSpec™ suite, IDS-AI™, IDS-NoC™, DFT automation)

Profile: Agnisys, Inc. — DAC 2026 press release, 23 July 2026

Demonstrated (as evidenced)	Asserted (company claim, unverified)
A DAC 2026 technical talk with named Google co-authors (Veerabhadrarao Vasa, Rajesh Gottumukkala of Google; Freddy Nunez of Agnisys) — the highest-evidence event in this profile [E2]	“Early adopters have seen up to 4× improvement in Test Coverage Per Flip-Flop (TCPF) and 50% fewer top-level ATPG patterns” — early adopters unnamed [E2]
Long-standing specification-automation business (IDesignSpec, SystemRDL/IP-XACT/PSS standards) [E2]	“correct-by-construction accuracy” — methodology claim, unproven in this release [E2]
Named DFT research direction: Design-Aware Test Points (DATP) in the SystemRDL-to-RTL flow [E2]	IDS-AI “drives the entire development flow from golden specification creation to RTL sign-off” — vendor framing [E2]

Agnisys — demonstrated vs asserted

Agnisys is the quietest profile here and the most instructive about evidence gradients: a two-decade spec-automation vendor whose headline numbers (4× TCPF, 50% fewer ATPG patterns) are vendor-reported like everyone else’s — but whose delivery vehicle is a DAC talk co-authored with Google engineers, one step from reviewable research [E2]. The gap between a talk and a published protocol with n, seeds and variance is exactly the gap this chapter exists to describe: the spec-to-RTL segment converges on agentic language, with measurement discipline still nobody’s differentiator.

CaretEDA

Category	Evidence tier	Funding anchor	Availability
Commercial open-stack (closest structural benchmark)	Press release [E2]	None stated in this release	“The 2026.08 release will be generally available in August 2026.” [E2]

Profile: CaretEDA, Inc. — DAC 2026 press release, 24 July 2026

Demonstrated (as evidenced)	Asserted (company claim, unverified)
Named tool base identical to SILICONA’s: “Simulation ... powered by Verilator”, “Logic Synthesis ... built on Yosys”, “Formal Verification ... built on SymbiYosys”, “Physical Synthesis ... powered by OpenROAD” [E2]	“the first time delivers a complete, commercially available Opensource EDA stack within the Spec-to-Netlist™ platform” — “complete” and “commercially available” are self-labels [E2]
A dated GA commitment (August 2026) and a concrete program: startup access plus “A 90-day evaluation license for CaretEDA’s agentic modules” [E2]	Agentic layer value: “pre-built, composable agents purpose-built for the rigor and complexity of chip design” (ABBs) [E2]
Microsoft booth demo slot at DAC 2026 (July 28); named CEO (Sashi Obilisetty) [E2]	Zero quantitative claims anywhere in the release — no customer names, tapeouts, or performance numbers [E2]

CaretEDA — demonstrated vs asserted

CaretEDA is the closest structural benchmark in the landscape: a company commercializing the same open stack SILICONA runs (Verilator, Yosys, SymbiYosys, OpenROAD) with an agentic layer on top — the identical thesis with the evidence layer inverted. Its release is honest about what it does not contain: not one performance number, not one customer name; the value proposition is packaging, support and a 90-day on-ramp. Even the company closest to SILICONA’s stack does not yet compete on published measurements; the differentiator is whether the product’s native output is a license or a traceable claim ledger.

OpenROAD

Category	Evidence tier	Funding anchor	Availability
Open-source flow (the open baseline)	Project docs x2 [E4] — strongest transparency tier	None (DARPA-origin project; no funding figures)	Fully open: BSD-3-Clause, 3.1k stars, 43,678 commits, active Sept 2026 [E4]

Profile: OpenROAD™ project — openroad.org + GitHub (commit 2026-09-11), live

Demonstrated (as evidenced)	Asserted (project framing)
Tapeout evidence, README-asserted and inspectable by construction: “over 600 silicon-ready tapeouts for technologies up to 12nm”; homepage: “1000+ runs and completed chip designs... from 180nm down to 12nm” [E4]	“RTL to GDSII in 24 hours with no human in the loop!” — capability framing, not a benchmark [E4]
Named institutional use: ETH Zurich 12nm RISC-V uLinux SoC; IIT Guwahati (3,000+ students); Tiny Tapeout (3,000+ first-time tapeouters); 500+ papers [E4]	“The leading open-source, foundational application for semiconductor digital design” — superlative, defensible by adoption [E4]
Live repo vitals observable at fetch: 3.1k stars, 43,678 commits, BSD-3; ships AI-agent contribution guidelines (.claude/, AGENTS.md) [E4]	“Autonomous, No-Human-In-Loop (NHIL) flow” — flow-level autonomy claim [E4]

OpenROAD — demonstrated vs asserted

OpenROAD is the only entity in this chapter whose adoption numbers are verifiable by construction, and it uses that position to make the same autonomy claims as the funded vendors — “no human in the loop” — but over an inspectable artifact. The tapeout counts (600+ silicon-ready, 12nm ceiling) are README-asserted, yet they sit inside an ecosystem that makes them the most checkable numbers in the commercial-EDA conversation. For the evidence-control-plane thesis, OpenROAD is both substrate and standard: SILICONA’s physical stages run on it, and its verifiability culture is the model this chapter’s counter-standard generalizes to AI-era claims.

LibreLane

Category	Evidence tier	Funding anchor	Availability
Open-source flow (community-maintained)	Project docs [E4]	None (FOSSi Foundation stewardship)	Fully open; plugin architecture, Nix reproducibility

Profile: LibreLane — librelane.readthedocs.io FAQ, live © 2020–2026

Demonstrated (as evidenced)	Asserted (project framing)
Governance history on record: “Since the liquidation of Efabless Corporation, the authors of OpenLane have lost control of the trademark, and OpenLane’s original developers currently maintain LibreLane under the stewardship of the FOSSi Foundation.” [E4]	“countless verified tapeouts” (OpenLane 1.0) — loose quantifier, anchored to Google MPW shuttles [E4]
Silicon proof tied to named shuttles: OpenLane 1.0 used for “more or less every open-source design on the Google MPW shuttles”; LibreLane/OL2 silicon-proven on “all Tiny Tapeout shuttles since 3.5” [E4]	Engineering-protocol posture: “Nix allows us to create a near-perfectly reproducible environment” [E4]
Published comparison matrix of the open-flow landscape (ORFS, SiliconCompiler, OpenLane, LibreLane) [E4]	Little hype; documentation prose throughout [E4]

LibreLane — demonstrated vs asserted

LibreLane closes the landscape with the chapter’s cleanest evidence behavior and its most instructive market fact. The evidence behavior: documentation-grade claims, named shuttles, and reproducibility treated as a first-class property (Nix, pinned environments) — the counter-standard already lives in open-source culture. The market fact: Efabless’s liquidation orphaned the OpenLane trademark and scattered stewardship to the

FOSSi Foundation — proof that open-EDA churn is real and that continuity of evidence (who can replay your tapeout in five years) is a survival feature. For the evidence-control-plane thesis, LibreLane is the precedent that the community values exactly what SILICONA systematizes.

5. Funding versus Substance

\$335M	\$131M	\$93M+	\$24.5M
Ricursive in ~4 months [E1]	ChipAgents Series A total [E1]	Cognichip total [S]	Agentrys total [E2]

The money is real and it validates the category. Within roughly four months of launch, Ricursive raised \$335M at a reported \$4B valuation [E1]; ChipAgents’ Series A total reached \$131M with Micron, MediaTek and Ericsson on the cap table [E1]; Cognichip announced a \$60M Series A bringing its total above \$93M [S]; Agentrys added \$24.5M led by Etna Labs with a MediaTek-led pre-seed [E2]. Add the incumbents’ agentic pushes [E2] and a semiconductor market whose base case runs from \$775B (2024) toward roughly \$1.6T by 2030 in the McKinsey-cited model [S], and the direction of the industry is not in doubt. Capital has decided that agentic chip design is a category.

But funding is not verified capability, and the journalism tier keeps saying so. TechCrunch’s Ricursive profile records zero shipped-product evidence and unnamed customers [E1]. Reuters’ ChipAgents exclusive quantifies the capital, the headcount and the investors — and no performance claim at all [E1]. The incumbent releases with the loudest multipliers are, on their own text, evaluation-stage: AMD “actively evaluating” [E2], capabilities “expected to be available to early-access customers in the second half of 2026” [E2], availability “planned for the second half of 2026” [E2]. This is not an accusation — pre-GA is a normal condition — but it means the market’s headline numbers are, today, mostly promissory.

Company	Raised	Tier	Shipped-product evidence on record here
Ricursive	\$335M @ \$4B	[E1]	None — no product details, benchmarks, customers or tapeouts in the journalism; customers unnamed
ChipAgents	\$131M Series A	[E1]	None quantified — CEO names verification as biggest speedup, unquantified; ~64 employees; strategic investors imply pilots
Cognichip	\$93M+ total	[S]	None — secondary tier only; no fresh primary source
Agentrys	\$24.5M	[E2]	Vendor-reported full-flow demo (32-bit CPU) + CVDP “exceed 90% accuracy” — checkable-in-principle, unverified; customers unnamed
Silimate	not published here	[E3]	Commercial product marketed; anonymous testimonials only; no named customers
Moore’sLabAI	not published here	[E3]	“Deployed at 15+ chip companies” self-reported; anonymous ROI figures; one named co-marketing partner (Breker)

Table 4: The funded cohort — capital (tiered) versus shipped-product evidence in this chapter’s sources.

The contrast SILICONA can honestly make

SILICONA’s counterweight to the funding narrative is a measured corpus in which every number is database-traceable: 563 recorded stage runs (81.0% PASS) across the 15-stage pipeline; 177 model evaluations in the AI ledger, of which 51 failed on provider rate limits (HTTP 429) and are published as failures, not smoothed away; and a demonstrated open-PDK closure at ladder level L9 — the sky130A reference design with DRC 0 violations, an LVS MATCH dual verdict (foundry deck plus deterministic structural compare) and GDS-II streamout, re-certified with post-route timing at revision 24 (950 mapped cells, 18,672.9 μm^2 , Fmax 240 MHz) [internal, measured]. Design-space search carries the same honesty: three families with measured improvements over baseline (+98.21% matrix-mac, +79.65% fir-filter across five runs spanning 74.15–83.32%, +34.30% stream-fifo on the multi-engine arena), one honest negative where a random control matched the LLM engines on a monotonic objective, and a failed grid family published as failed. Model neutrality is architectural, not measured performance: a 12-provider registry is implemented, and exactly one provider (GLM) is measured to date — the 429 failures are the visible proof of that single-provider dependency. None of this claims SILICONA out-agents anyone; it claims that SILICONA’s numbers can be checked, and that the checking is the product.

6. Positioning: Where the Control Plane Sits

The matrix below states how SILICONA differs from the landscape without ranking anyone. Each row is a design commitment of the SILICONA platform, checked against the pattern documented in Chapters 2–5 — the comparison is between an engineering posture and a market pattern, not between two products’ benchmark scores, because no competitor publishes scores that could be compared.

Dimension	SILICONA (as built)	Landscape pattern (as documented)
Evidence discipline	Per-run provenance on every result (revision id, provider+model, tool versions); a consistency audit that blocks non-conforming reports	Headline multipliers with no protocol, n, seeds or variance ([E2]/[E3] profiles, Table 3)
Traceable workflows	Engineering graph: every artifact hash-bound; every AI change a new auditable revision; failures classified with root causes (104 recorded)	Company-reported outcomes at best; anonymous testimonials at worst; no per-run publication
Model neutrality	12-provider registry behind three wire protocols — an architectural claim; exactly one provider measured (GLM, 177 evals incl. 51 honest 429 failures)	Model partnerships as credibility (NVIDIA Nemotron/OpenShell across incumbents [E2]); no provider-neutral measurement story anywhere
Measured vs asserted	Three enforced evidence states (MEASURED / ANALYTICAL ESTIMATE / CONFLICTING EVIDENCE); power stays labeled ESTIMATE; demonstrated open-PDK closure claimed only as demonstrated	Autonomy-level taxonomies (“Level-5” [E2]); “sign-off-clean” and “no human in the loop” claims without published runs [E2]/[E3]
Deployment honesty	Published gaps: 429 failure rows, failed search families, single-provider measurement, sandbox-bound deployment notes	Pre-GA framing present but self-reported; failures, when mentioned, are anecdotes not ledgers

Table 5: SILICONA’s posture versus the documented landscape pattern (no competitor is ranked).

Where SILICONA does not compete

- **Better models.** No claim that SILICONA’s models beat anyone’s — model neutrality is the architecture (a 12-provider registry), not a measured result; one provider measured to date.
- **Faster autonomy multipliers.** No iteration-time or speedup claims against humans or competitors — the standing do-not-publish rules from the white paper claim ledger.
- **Foundry sign-off.** The claim ceiling is demonstrated open PDK closure on sky130A: DRC 0, LVS MATCH dual verdict, GDS-II streamout. Post-route extraction and signoff STA are explicitly unclaimed; “sign-off” appears in this document only inside competitors’ quotations.
- **Production readiness.** The platform’s own published gaps stay published; nothing here asserts deployment-readiness beyond what the corpus shows.

That leaves the position this chapter has been circling: SILICONA does not compete on autonomy theater, model prowess or headline speedups — the funded and incumbents own and defend those positions well. It competes on the thing none of them publish: per-run, replayable, failure-inclusive evidence. Where every competitor’s numbers are unverifiable by construction, the only verifiable shop is differentiated by default.

7. Implications

Three implications close this chapter. First, the market gap is real and oddly uncontested: across fourteen primary sources and four evidence tiers, not one commercial vendor leads with verifiable evidence discipline — claims concentrate where evidence is thinnest (autonomy multipliers [E2][E3]), and the only checkable numbers come from outside the commercial layer (OpenROAD’s inspectable record [E4]; Agentrys’ CVDP number [E2]). An evidence-driven engineering control plane sits in that gap as infrastructure: not the loudest agent, but the ledger every agent’s claims would have to survive.

Second, the funded cohort is SILICONA’s future customer set as much as its competition: a market that has priced learning-across-designs stories at billions of dollars [E1] will eventually need the machinery that makes such stories checkable — the very machinery SILICONA has already built and measured on itself. Third, the standing offer: the companion SILICONA Benchmark Protocol (v1.0) is published alongside this chapter as a comparable-claims standard any vendor may adopt. If the market adopts it, claims become comparable and SILICONA’s discipline becomes the baseline; if it does not, SILICONA remains the only player whose numbers anyone can check. Either outcome is a good position; the protocol is the offer that forces the choice.

8. Source Register

Fourteen primary sources were fetched live for this edition on 12 September 2026 (02:30–02:32 UTC); fourteen of fourteen returned HTTP 200 with extractable body content — zero failures, zero retries. Raw payloads and fact sheets are on file in the SILICONA research store. Secondary foundation sources are listed separately and labelled: not freshly fetched, context only, never the basis of a SILICONA contrast.

Publisher / source	Date	URL	Tier
Synopsys newsroom (PRNewswire)	2026-07-27	https://news.synopsys.com/2026-07-27-Synopsys-Advances-Agentive-AI-Chip-Design-with-A-MD-and-Microsoft	[E2]
Synopsys newsroom (PRNewswire)	2026-07-26	https://news.synopsys.com/2026-07-26-Synopsys-Showcases-Comprehensive-Autonomous-Engineering-Workflows-from-Silicon-to-Systems%2C-Developed-with-NVIDIA-Technology	[E2]
Cadence newsroom	2026-06-01	https://www.cadence.com/en_US/home/company/newsroom/press-releases/pr/2026/cadence-unveils-industrys-first-fully-autonomous-virtual.html	[E2]
TechCrunch (Julie Bort)	2026-02-16	https://techcrunch.com/2026/02/16/how-recursive-intelligence-raised-335m-at-a-4b-valuation-in-4-months/	[E1]
Reuters (Max A. Cherney)	2026-07-29	https://www.reuters.com/business/nvidia-partner-chipagents-raises-60-million-accelerate-chip-design-with-ai-2026-07-29/	[E1]
Agentrys newsroom	2026-08-26	https://agentrys.ai/news/agentrys-raises-24-5-million	[E2]
silimate.com (homepage)	live © 2026	https://www.silimate.com/	[E3]
moorelab.ai (homepage + FAQ)	live © 2026	https://www.moorelab.ai/	[E3]
openroad.org (project homepage)	live	https://openroad.org/	[E4]
GitHub — The-OpenROAD-Project/OpenROAD	live (commit 2026-09-11)	https://github.com/The-OpenROAD-Project/OpenROAD	[E4]
caretada.com newsroom	2026-07-24	https://caretada.com/news/dac-2026/	[E2]
rise-da.com (company page)	live	https://www.rise-da.com/company/	[E3]
agnisys.com newsroom	2026-07-23	https://www.agnisys.com/news/agnisys-to-showcase-agentive-ai-network-on-chip-and-dft-automation-breakthroughs-at-dac-2026/	[E2]
librelane.readthedocs.io (FAQ)	live © 2020–2026	https://librelane.readthedocs.io/en/latest/faq.html	[E4]

Table 6: Primary sources — publisher, date, URL, evidence tier.

#	Source	Basis	Used for
S1	Siemens EDA dossier [7][8]	Foundation extract	Siemens EDA profile (EDA AI System; Solido)
S2	Cognichip dossier [11][12]	Foundation extract	Cognichip profile (\$60M Series A; \$93M+)
S3	McKinsey-cited economics [22][23][24][25]	Foundation extract	Category context only (\$775B → ~\$1.6T)

Table 7: Secondary foundation sources — context-only, not freshly fetched.

Fetch discipline for the record: URLs were captured utm-stripped with timestamps and raw payloads retained; quotes were copied verbatim; where a source contradicts itself (Synopsys’s 40% vs 25–40%), both readings are quoted. Corrections and retractions follow the Benchmark Protocol: recorded, dated, never silent.